

User-level Grid Monitoring with Inca 2

Shava Smallen, Kate Ericson, Jim Hayes, Catherine Olschanowsky
San Diego Supercomputer Center
University of California, San Diego
9500 Gilman Drive, La Jolla, CA 92093-0505, USA
{ssmallen,kericson,jhayes,cmills}@sdsc.edu

Abstract

User-level Grid monitoring is valuable, how it relates to other Grid monitoring, and what are its implementation challenges. Inca 1 good but had limitations. In this paper, we introduce Inca 2 and describe its features and architecture. We then show a few use cases for Inca is being deployed on TeraGrid and XXX? System impact results and performance results. Finally, we discuss our future work.

1. Introduction

Grid systems provide unified and coherent access to distributed computing, data storage and analysis, instruments, and other resources. These systems require the careful coordination of software packages, services, and configurations across multiple, heterogeneous resources. The TeraGrid project [12], for example, manages the coordination of software and services by deploying and monitoring a common user environment across distributed, heterogeneous resources. The organization of TeraGrid software and services is referred to as the Common TeraGrid Software and Services (CTSS). TeraGrid's software and services uniformity simplifies access to its resources, which consist of more than 102 teraflops of compute capability and more than 15 petabytes of online storage, all interconnected by a network that can transfer a terabyte of data in under 10 minutes.

Providing and maintaining a stable infrastructure for these complex Grid systems poses challenges for both administrators who build and maintain Grid resources and scientists who use them. First, system administration is distributed across multiple administrative domains which requires a significant amount of coordination. This is typically performed by a group of Grid operators who take into account local site policies and decide when software (including updates and patches) should be deployed to resources and how it should be configured for the set of het-

erogeneous resources. Each site's system administrators, who are oftentimes not Grid experts, are then responsible for providing the required Grid software components and services for their resource and debugging any problems that arise. Well-defined requirements and good communication are required in this model. Otherwise inconsistencies between the resources arise that can either inconvenience the user or prevent them from using particular resources together. Second, as a large distributed system, failures will occur over time in the Grid services due to network or system failures, software misconfiguration, or software bugs. In order to detect such problems, monitoring of the resources is required.

One approach to monitoring Grids, used by tools such as MonALISA [6] and GridICE [1], is to aggregate and display data from existing cluster or system administrator monitoring tools such as Ganglia [8], CluMon [3], and Nagios [10]. This provides a centralized, systems-level view of Grid resources where low-level host statistics and queue information can be examined. This type of monitoring information is useful for showing the utilization of Grid resources, however, it does not provide the type of high-level monitoring needed to detect user problems within the Grid infrastructure, such as incompatible software versions.

User-level Grid monitoring approaches testing and performance measurement of the Grid infrastructure from the user perspective. The goal is to act as an impartial test user of the Grid and detect problems with the infrastructure so they can be fixed before users notice them – user complaints should not be the first indication of Grid failures. We refer to an executable program that tests or measures some aspect of a system or installed software from the user-perspective as a reporter. In our view, the following guidelines define user-level Grid monitoring:

- In order to reflect regular user experiences, tests or performance measurements of the Grid infrastructure should be done from a standard user account. In fact, it is important to not execute under a system admin-

istrator's account because they may be privileged and often have custom shell initialization files.

- Also in order to reflect regular user experiences, tests should be written and configured using information directly from user documentation (e.g., hostnames, ports, pathnames, etc.). This may not always be possible when documentation and testing are done in parallel in a pre-production environment but they should be closely coordinated activities.
- Since Grids are dynamic, user-level tests or performance measurements should be automated and executed periodically.
- Tests or performance measurements that interact with Grid services require authentication and should be executed using a standard user GSI credential that is mapped to the standard user account.
- Tests or performance measurements should be executed locally on all Grid resources when appropriate so that all Grid access points available to users are verified. Similarly, it is important to execute some tests all-to-all in order to detect site-to-site configuration errors such as authentication problems.

In our experience, writing and maintaining user-level tests and performance measurements is an iterative refinement process since Grids are dynamic and the software environment changes over time as packages are upgraded. It is also difficult to write and deploy a test perfectly the first time because of portability issues and sometimes incomplete user documentation. For example, it may take a few iterations to determine whether detected failures are the result of a faulty test, faulty user documentation, or a faulty system. Realistically, this is a labor-intensive activity and thus the amount of user-level Grid monitoring is often limited by the amount of people resources that can be assigned to it. However, a good goal is to test all the commands that exist in the user documentation and/or have at least one basic test for each software component of the Grid infrastructure.

In 2003, SDSC, in partnership with TeraGrid, began developing Inca 1 [11] to implement a user-level monitoring system with the above features and started using it to validate and verify that CTSS was deployed consistently across all TeraGrid resources and to monitor its status. At that time, the only available user-level Grid monitoring tools were the NCSA TestGrid script [9] and GITS [5], scripts that run a fixed number of Grid tests and output results in HTML. Although these tools were easy to install and produced useful information, they showed only the view of the Grid from a single resource, lacked automation, and were not easily extensible. The initial version of Inca was implemented as a client-server architecture. It provided data

collection, data storage (with limited archiving capabilities) that was accessible from a Web services interface, and data display through Web status pages. Inca was first deployed to TeraGrid in mid 2003 and released in late 2003. After running Inca 1 for a year and a half on TeraGrid, we learned some valuable lessons and designed Inca 2, our current release. Inca 2 contains a number of substantial improvements over Inca 1 with respect to security, installation and maintenance, and storage capabilities. A production release of Inca 2 was provided in February 2007 and has been running on TeraGrid since November 2006. Currently, seven other Grids in the U.S., Europe, and Australia use Inca.

In the next section, we describe the features of Inca 2 and note which are new features based on our experiences with deploying Inca 1 on TeraGrid. We then describe the Inca 2 architecture and how it can be used to provide user-level Grid monitoring. In Section 4, we describe two uses of Inca 2 for Grid software environment verification and benchmarking. Finally, we describe some future work and summarize the paper.

2. Inca 2 Features

The goal of the Inca 2 design was to

1. A reporter
2. logging in reporters
3. Given the challenges of writing portable tests and benchmarks, Inca 2 includes a set of Perl APIs that provide convenience functions to simplify this process.
4. Periodic tests executed by the monitoring system should be automated and run from a standard user account. A user GSI credential should be issued and used by the monitoring system to test services that require authentication.
5. Tests should be executed locally on all Grid resources when appropriate so that all Grid access points available to users are tested.
6. The configuration of tests should be managed centrally in order to ensure consistent testing of the resources, rather than by system administrators who may not use the same information that is available to users.
7. The testing system should easily adapt to new resources and testing requirements. In our experience, testing has an iterative refinement process because Grids are dynamic and the software environment changes over time as packages are upgraded. Furthermore, it is often difficult to write and deploy a test perfectly the first time because of portability issues and incomplete user documentation.

8. System administrators should not need to know the details of the monitoring system in order to troubleshoot problems. When a problem is detected the monitoring system should provide system administrators with sufficient information to locate the problem easily and test repairs.
9. A history of test results (especially any error messages) should be retained in order to understand the behavior of their Grid over time. This history should be available through a flexible querying interface.
10. A running monitoring system should have a limited impact on the resources it measures and should be secure.

3. Inca 2 Architecture

Figure 1 shows the architecture of Inca 2 designed to implement the features described in the previous section. Inca 2 is composed of 3 core components (in highlighted box): the agent, depot, and reporter manager. The *agent* and *reporter managers* coordinate the execution of tests and performance measurements on the Grid resources and the *depot* stores and archives the results. The inputs to Inca 2 are one or more *reporter repositories* that contain user-level tests and benchmarks called *reporters* and a configuration file describing how to execute them on the Grid resources created using an administration GUI tool called *incat* (Inca Administration Tool). The output or results collected from the resources are queried by the *data consumer* and displayed to users. The following steps describe how an Inca administrator would deploy user-level test and/or performance measurements to their resources.

1. Using the guidelines described in the introduction, the Inca administrator either writes reporters to monitor their Grid or uses existing reporters in a published repository.
2. The Inca administrator creates a deployment configuration file that describes the user-level monitoring for their Grid using *incat* and submits it to the agent.
3. The agent
 - (a) fetches reporters from the reporter repository
 - (b) creates a reporter manager on each resource
 - (c) sends the reporters and instruction for executed them to each reporter manager.
4. Each reporter manager executes reporters according to its schedule and sends results (reports) to the depot
5. Data consumers display collected data (reports) by querying the depot

The following subsections describe the Inca components in more detail using the order of the steps above.

3.1. Reporters

An Inca reporter is an executable program that tests or measures some aspect of a system or installed software. Reporter executables are designed to be easy to produce and can be run outside of the Inca system (e.g., by a system administrator). A reporter should be written following the user-level monitoring guidelines described in the introduction and could be a simple Globus [4] GRAM gatekeeper ping test (see Figure 2) or a more complex Grid application benchmark [2]. Reporters must support certain command line options and produce an XML report according to the Inca reporter schema. The goal of the reporter schema is to accommodate multiple types of data and to capture information about the reporter execution to diagnose any detected failures. The reporter schema consists of the following elements:

- a set of header tags describing the context of the reporter execution: GMT timestamp, hostname, reporter name, reporter version, working directory, reporter path, and arguments;
- an optional set of debug or information log tags similar to log4j [7] output;
- a body containing the results expressed as any XML sequence; and
- an exit status tag indicating whether the reporter was able to complete its test or measurement and an optional error message.

Because the body of the report can be any XML sequence, it enables reporters to express a wide variety of information. Today, we have four standard body schemas to express software version information, software functionality or service tests results, usage information, and performance results.

We provide an extensible set of Perl APIs to handle much of the effort of writing reporters. Figure 2 shows (a) an Inca reporter that pings a pre-WS Globus GRAM server and (b) its corresponding output when executed on a local machine. The reporter uses two of the Perl APIs: `Inca::Reporter::SimpleUnit` (which is a subclass of `Inca::Reporter`) and `Inca::Reporter::GridProxy`. The `Inca::Reporter` API provides convenience functions for handling the required input arguments (lines 2(a) 12-14 → lines 2(b) 10-31) and printing the Inca-compliant header and exit status XML tags (lines 2(a) 4-11, 24 → 2(b) 1-9, 43-46). It also includes some informational and debug logging functions again similar to log4j; the log output has been very useful to system administrators who can view a summary

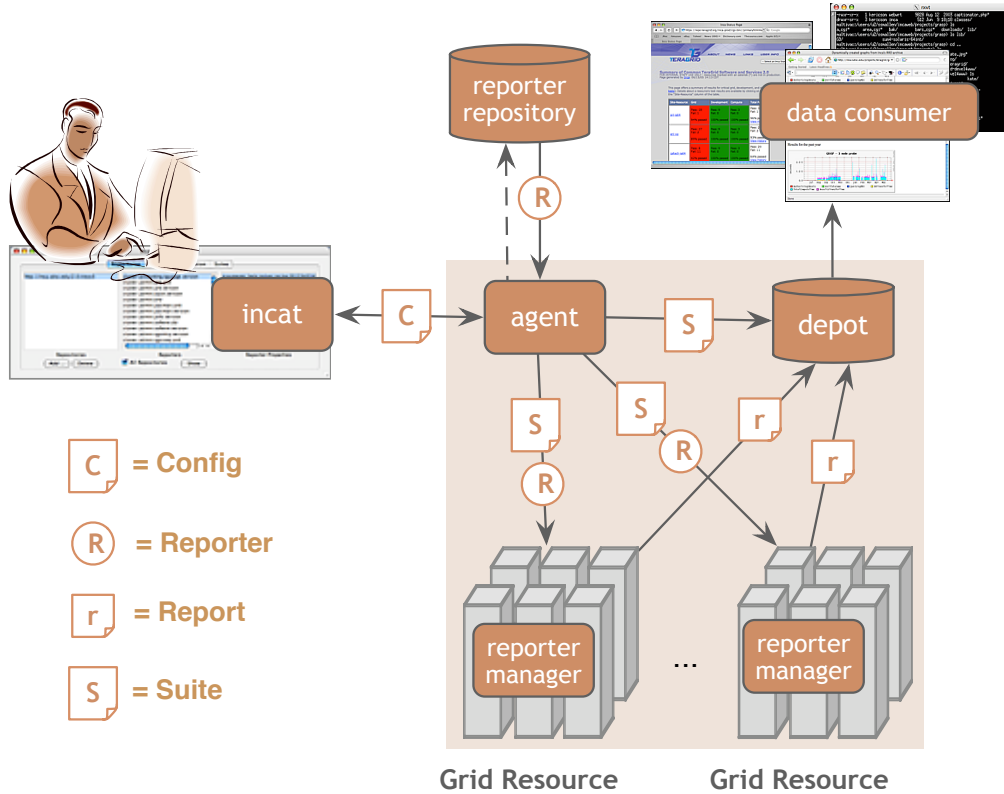


Figure 1. Inca architecture.

of the test or benchmark that the reporter performed without reading the reporter source code. Line 2(a) 15 shows a special log function, `loggedCommand`, that executes a system command (using a 30 second timeout) and logs to the 'system' level of the log XML output (lines 2(b) 32-36). The subclass `Inca::Reporter::SimpleUnit` is the API responsible for printing the simplest standard reporter body schema – reporting software functionality or service tests results; it handles the printing of the small body XML (lines 2(b) 39-41) using the functions `unitSuccess` (line 2(a) 22) or `unitFailure` (lines 2(a) 18, 20) depending on the output of system command. And finally, the `Inca::Reporter::GridProxy` API (line 2(a) 3) adds a proxy dependency to the reporter which tells the Inca system to download a short-term proxy before running this reporter (this is discussed further in Section ??). Most current reporters use the Perl APIs and consist of fewer than 30 lines of code.

3.2. Reporter Repositories

The Inca agent retrieves reporters from collections called reporter repositories, that consist of reporters, required

packages and libraries, and a catalog file. Repository contents are accessible using a URL and can be shared across multiple Inca deployments. The Inca team publishes a reporter repository that contains reporters developed for TerraGrid.

3.3. Inca administration GUI tool (incat)

The Inca administrative GUI tool (incat) is the main interface by which an Inca administrator interacts with the Inca system. Incat allows the Inca administrator to view the choose which resources to monitor and which tests or benchmarks to deploy to those resources. The following subsections describe the Inca components that are configurable via incat. To configure reporter execution, an Inca administrator launches a GUI tool, called incat. Incat allows the administrator to choose which resources to monitor and which reporters to deploy to those resources. For each reporter, the administrator can specify the resources to run on, command-line arguments, the runtime environment, the frequency of execution, and resource limits. The Inca administrator may also group related reporters into suites that

```

#!/usr/bin/env perl
2 use Inca::Reporter::SimpleUnit;
3 use Inca::Reporter::GridProxy;
4 my $reporter = new Inca::Reporter::SimpleUnit(
    name => 'grid.globus.gramPing',
    version => 2,
    description =>
6     'Checks gatekeeper is accessible from local machine',
    url => 'http://www.globus.org',
    unit_name => 'gramPing'
10 );
12 $reporter->addArg('host', 'gatekeeper host');
13 $reporter->processArgv(@ARGV);
14 my $host = $reporter->argValue('host');
15 my $out = $reporter->loggedCommand
    ("globusrun -a -r $host", 30);
16 if (!$out) {
17     $reporter->unitFailure("globusrun failed: $!");
18 } elsif($out =~ /GRAM Authentication test successful/) {
19     $reporter->unitFailure("globusrun failed: $out");
20 } else {
21     $reporter->unitSuccess();
22 }
23 $reporter->print();

```

(a)

```

<?xml version='1.0'?>
2 <rep:report
    xmlns:rep='http://inca.sdsc.edu/dataModel/report_2.1'>
3   <gmt>2007-02-22T18:05:37Z</gmt>
4   <hostname>client64-51.sdsc.edu</hostname>
5   <name>grid.globus.gramPing</name>
6   <version>2</version>
7   <workingDir>/Users/ssmallen</workingDir>
8   <reporterPath>grid.globus.gramPing-2</reporterPath>
9   <args>
10     <arg>
11       <name>help</name>
12       <value>no</value>
13     </arg>
14     <arg>
15       <name>host</name>
16       <value>localhost</value>
17     </arg>
18     <arg>
19       <name>log</name>
20       <value>3</value>
21     </arg>
22     <arg>
23       <name>verbose</name>
24       <value>1</value>
25     </arg>
26     <arg>
27       <name>version</name>
28       <value>no</value>
29     </arg>
30   </args>
31   <log>
32     <system>
33       <gmt>2007-02-22T18:05:36Z</gmt>
34       <message>globusrun -a -r localhost</message>
35     </system>
36   </log>
37   <body>
38     <unitTest>
39       <ID>gramPing</ID>
40     </unitTest>
41   </body>
42   <exitStatus>
43     <completed>true</completed>
44   </exitStatus>
45 </rep:report>

```

(b)

Figure 2. (a) An Inca reporter that tests the user-level availability of a Globus GRAM server and (b) its sample output.

can be shared across different Inca deployments. This can be useful to determine interoperability among Grids or to determine whether an application's requirements are being fulfilled on a Grid.

3.3.1 Suites

3.3.2 Resource Macros

3.4. Agent

The Inca agent is a server that implements the configuration specified by the Inca administrator. After determining which Inca tests should be executed on each resource, the agent stores this configuration information in the depot. It then stages and launches a dependent component, called a reporter manager, on each resource using either SSH or Globus [13]. Once a reporter manager contacts its agent, the agent transmits Inca tests to execute, along with their dependencies, configuration and schedule of execution.

3.5. Reporter Manager

The Inca reporter manager is a lightweight process responsible for managing the schedule and execution of Inca tests, called reporters, on a single resource. The reporter manager receives reporter updates and dependencies (e.g., Inca Perl APIs or source code) from the agent along with requests for reporter scheduling changes. Running under a regular user account, the reporter manager executes reporters on-demand or using an internal cron scheduler, and sends reports to the depot for archiving. The reporter manager monitors reporter system usage and enforces limits (e.g., wall clock time, CPU time, memory). System usage information is sent to the depot with each report.

3.6. Depot

The Inca depot server is responsible for storing configuration information and the data produced by reporters. The depot maintains a relational database via Hibernate [14] so that it can use a variety of databases. The depot provides full archiving of reporter output and structures its schema to reduce redundant data. Data can be queried using SQL queries. Predefined queries exist to return the latest report instances of a suite, a single report instance, or a report history. A Web services interface is also available to provide unauthenticated query access to data.

3.7. Data Consumer

The Inca data consumer is a Web application that queries the Inca depot for data and displays it in a user-friendly format. The data consumer is packaged with Jetty [15] so that

pages can be served immediately without deploying Tomcat [16]. A set of JSP tags and pages query the Inca depot for data (returned as XML) and a set of XSL stylesheets format the data as HTML.

4. Use Cases

4.1. TeraGrid

4.1.1 Description and Requirements

4.1.2 Usage of Inca

4.1.3 Results

4.2. GrASP

4.2.1 Description and Requirements

4.2.2 Usage of Inca

4.2.3 Results

4.3. Other uses

5. Future Work

5.1. Graphing and Summary Statistics

5.2. Knowledge Base

5.3. Fault Tolerance

6. Summary

References

- [1] S. Andreatozzi, N. D. Bortoli, S. Fantinel, A. Ghiselli, G. L. Rubini, G. Tortone, and M. C. Vistoli. GridICE: a monitoring service for Grid systems. *Future Generation Computer Systems*, 21(4):559–571, 2005.
- [2] G. Chun, H. Dail, H. Casanova, and A. Snavely. Benchmark Probes for Grid Assessment. April 2004.
- [3] Clumon Cluster Monitoring Web page. <http://clumon.ncsa.uiuc.edu>, 2007.
- [4] I. Foster and C. Kesselman. Globus: A Metacomputing Infrastructure Toolkit. *International Journal of Supercomputer Applications*, 11(2):115–128, 1997.
- [5] The UK Grid Integration Test Script - GITS. <http://www.soton.ac.uk/~djb1/gits.html>, 2007.
- [6] I. Legrand, H. Newman, R. Voicu, C. Cirstoiu, C. Grigoras, M. Toarta, and C. Dobre. MonALISA: An Agent based, Dynamic Service System to Monitor, Control and Optimize Grid based Applications. In *CHEP 2004*, Interlaken, Switzerland, September 2004.
- [7] Log4J Web page. <http://logging.apache.org/log4j>, 2007.
- [8] M. Massie, B. Chun, and D. Culler. The Ganglia Distributed Monitoring System: Design, Implementation, and Experience. *Parallel Computing*, April 2004.
- [9] NCSA TestGrid Project. <http://grid.ncsa.uiuc.edu/test>, 2007.
- [10] Nagios Web page. <http://www.nagios.org>, 2007.
- [11] S. Smallen, C. Olschanowsky, K. Ericson, P. Beckman, and J. M. Schopf. The Inca Test Harness and Reporting Framework. In *Proceedings of Supercomputing 04*, November 2004.
- [12] The TeraGrid Project Web page. <http://www.teragrid.org>, 2007.