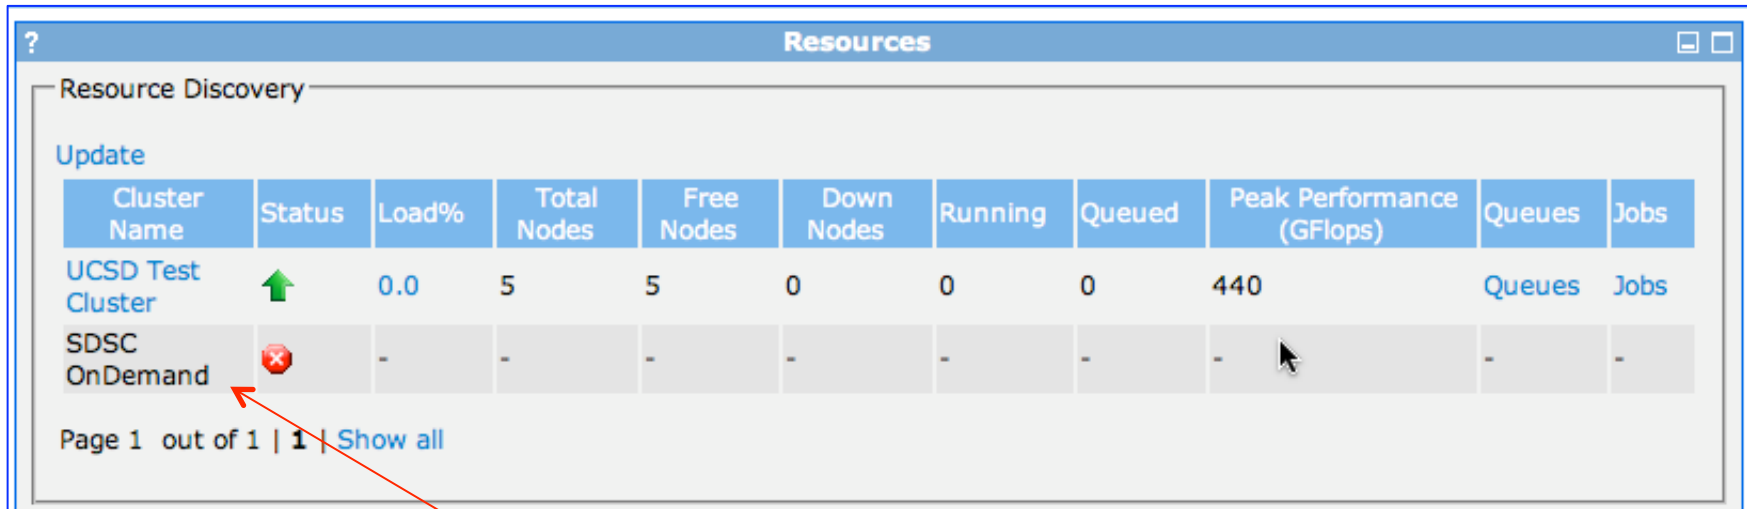

Inca User-level Grid Monitoring

Shava Smallen

ssmallen@sdsc.edu

Things fail ...



Cluster Name	Status	Load%	Total Nodes	Free Nodes	Down Nodes	Running	Queued	Peak Performance (GFlops)	Queues	Jobs
UCSD Test Cluster	↑	0.0	5	5	0	0	0	440	Queues	Jobs
SDSC OnDemand	✗	-	-	-	-	-	-	-	-	-

Page 1 out of 1 | 1 | [Show all](#)

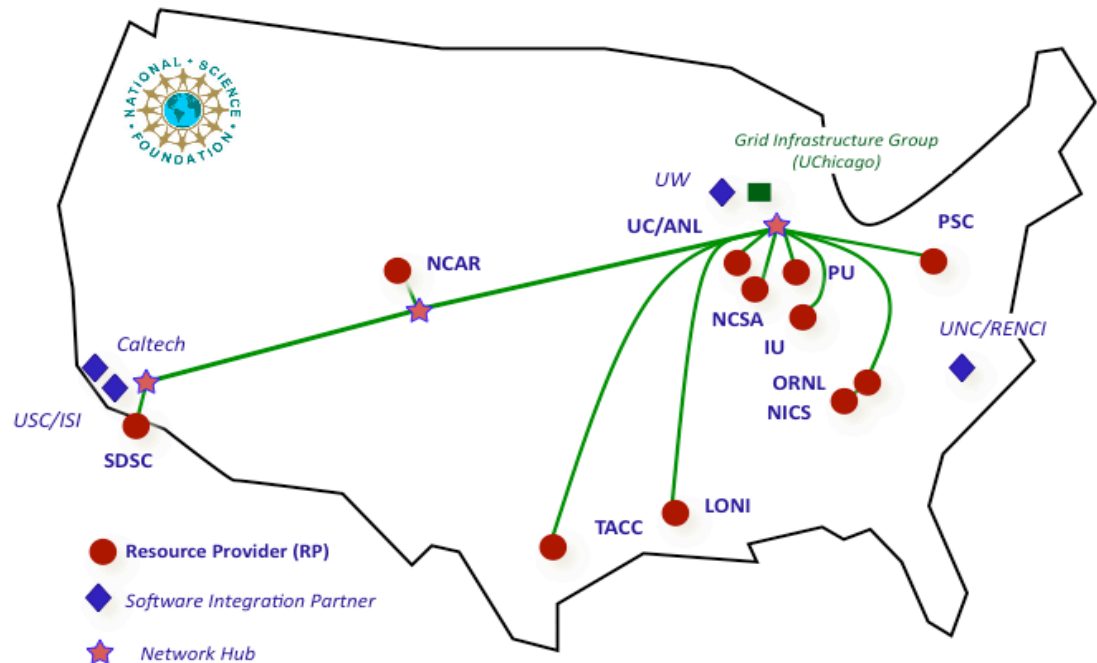
For example,

- Resource does not display in portal (MDS not running)
- Cannot submit job (WS-GRAM not running)

Distributed System: *“You know you have one when the crash of a computer you've never heard of stops you from, getting any work done.”*
-Leslie Lamport

Goal: reliable grid software and services for users

- Grids provide unified access to resources that advance scientific exploration
- Unfortunately, complex to deploy and maintain
- Detect and analyze failures




TeraGrid™

*1+ PFlop compute power,
30+ PB of online and archival data storage,
multi-Gpbs links*

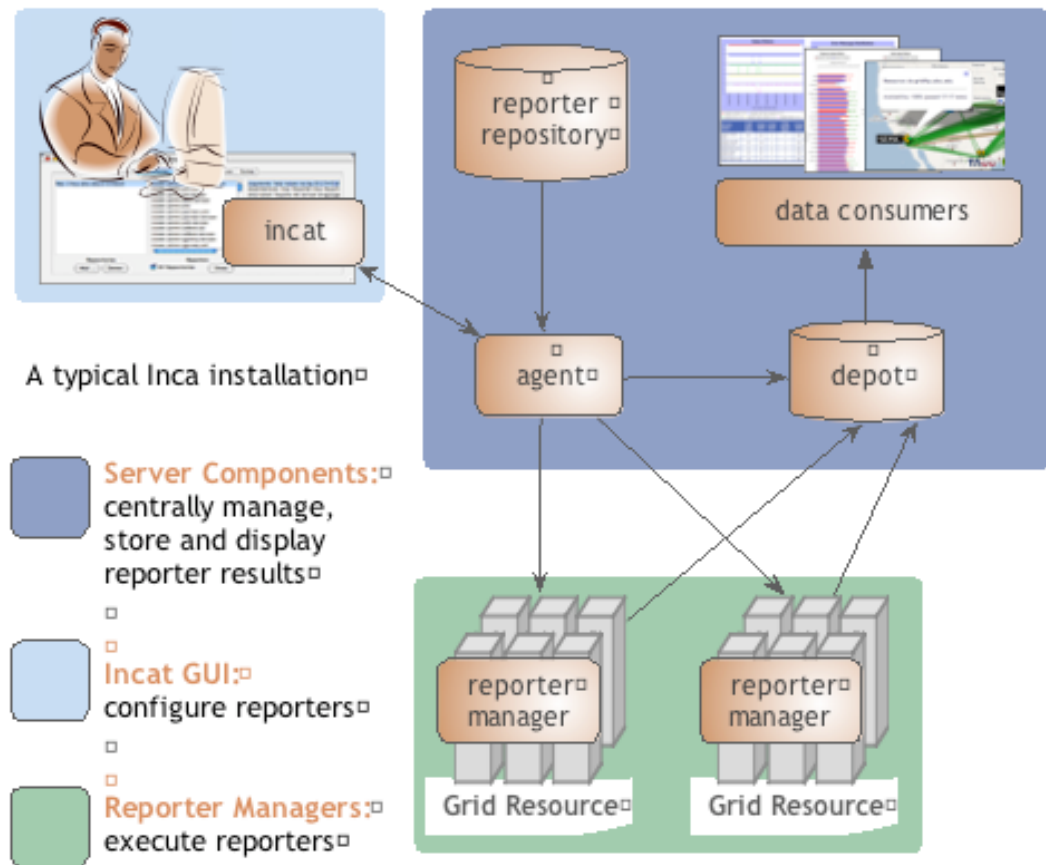
Who benefits from grid monitoring?

- Grid managers
 - Verify requirements are fulfilled by resource providers
 - Identify failure trends
- System administrators
 - Email notification
 - Debugging support
- End users
 - Higher availability
 - Debug user account/environment issues



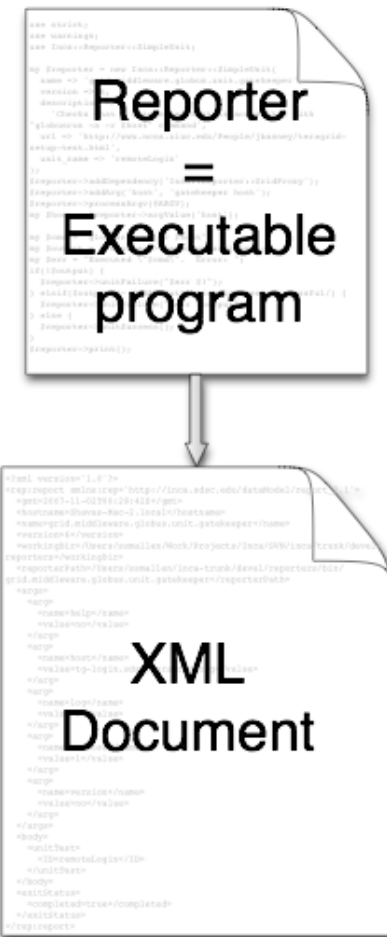
Inca provides user-level grid monitoring

- Enables consistent user-level testing across resources
- Easy to configure and maintain
- Easy to collect data from resource
- Archived results support troubleshooting
- Large variety of tests
- Comprehensive views of data

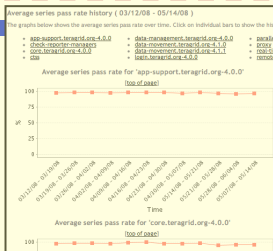


Reporters collect monitoring data

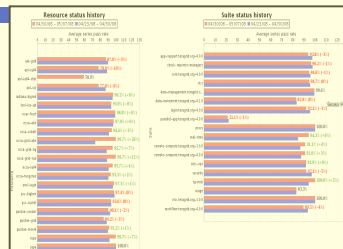
- Executable programs that measure some aspect of the system or installed software
- Supports a set of command-line options and writes XML to stdout
- Schema supports multiple types of data
- Extensive library support for perl and python scripts (most reporters < 30 lines of code)
- Independent of other Inca components



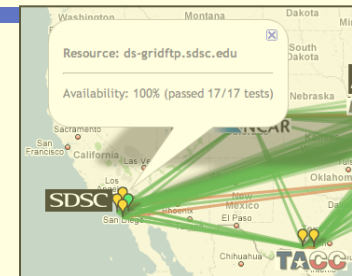
Tests Summary



Resource status history




Weekly status report



Cumulative test status by resource


Error history summary

Error history summary



Related test histories

Average series pass rate for 'data-movement.teragrid.org-4.1.0'



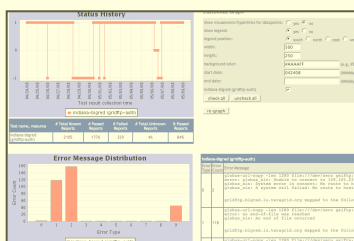
Average series pass rate for 'data-movement.teragrid.org-4.1.2'



gsi-openssh	resource1	resource2	resource3	resource4
version: >= 3.9	4.6p1	4.6p1	4.5p1	4.5p1
gsi-sh-unit	pass	pass	error	pass
gx-map	resource1	resource2	resource3	resource4
version: 0.5.3.3(0.5.3.2p1)	0.5.3.2p1	0.5.3.2p1	0.5.3.2p1	0.5.3.2p1
myproxy	resource1	resource2	resource3	resource4
version: >= 3.4	3.4	3.4	3.4	3.4
softnrv	resource1	resource2	resource3	resource4
version: 1.6.2	1.6.2	1.6.2	1.6.2	1.6.2
softnrv-unit	pass	pass	pass	pass
tgproxy	resource1	resource2	resource3	resource4
tgproxy-unit	pass	pass	pass	pass

Test status by package and resource


Test Details



Individual test history

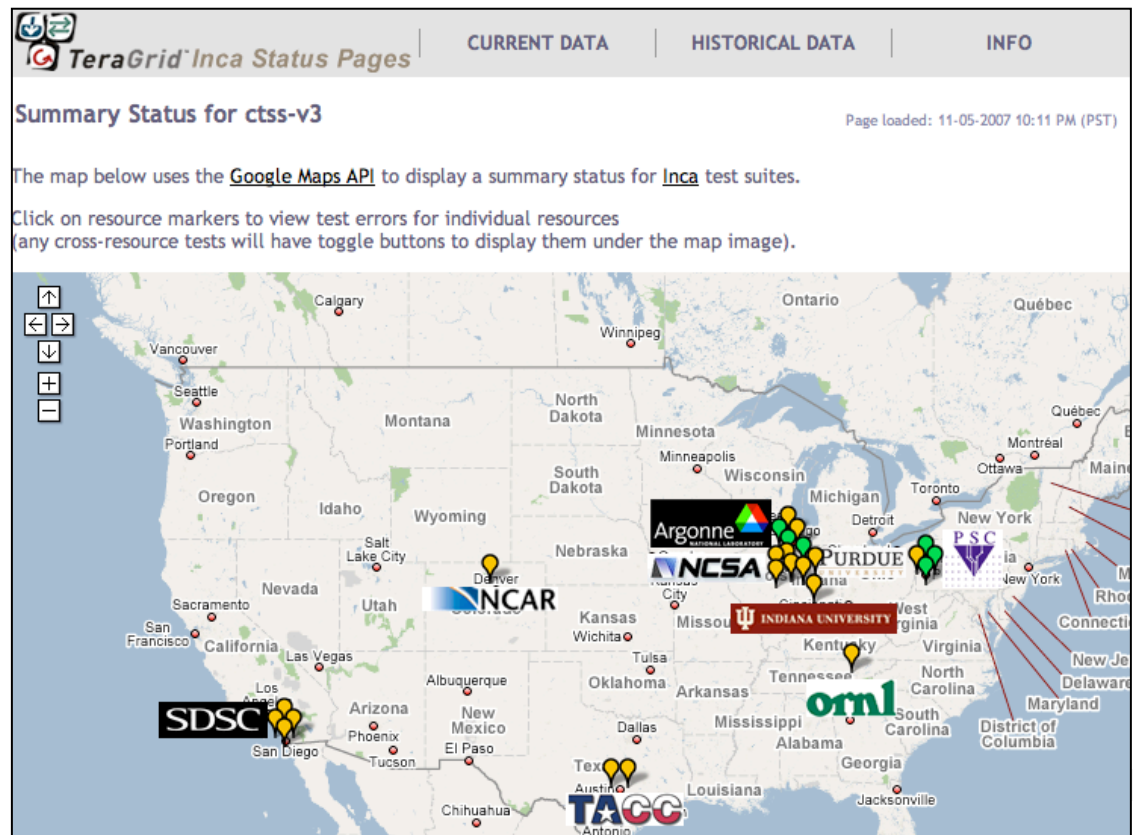
Details for spu-vgst series	
Completed	view result history
Reporter details	
reporter name	data-access-unit-test-suite
reporter version	4
Execution information	
run at	05-08-2008 07:45 AM (PDT)
age	13 hours 26 min
run on (hostname)	lg-login1.sdsc.teragrid.org
memory usage (MB)	13,8594
cpu time (secs)	1.49512
build clock time (secs)	0.48403
Input parameters	
help	no
log	5
verbose	1
debug	0

Individual test result details



Inca TeraGrid deployment

- Running since 2003
- Total of 2 tests running on 20 login nodes, 3 grid nodes, and 3 servers
- Coordinated software and services
- Cross-site tests
- GRAM usage
- CA certificate and CRL checking
- Resource registration in information services



Screenshot of Inca status pages for TeraGrid

<http://inca.teragrid.org/>

Inca GEON deployment

- Running since Feb 2008
- Total of 206 tests running on 5 login nodes and 6 servers
- LiDAR workflow services
- Web servers
- Ssh connectivity
- Base system information (Rocks, Gcc, Java, etc.)



Screenshot of Inca status map for GEON

<http://inca-geon.sdsc.edu>

Monitoring UCSD's campus portal

- At first glance: (remote tests)
 - Tomcat/Gridsphere
 - MyProxy
 - MDS
 - WS-GRAM

CampusPortal	
<ul style="list-style-type: none">• grid.ucsd.edu_homePage• grid-test.ucsd.edu_homePage• myproxyPing	
ucsd-portal	
grid.ucsd.edu_homePage	pass
grid-test.ucsd.edu_homePage	pass
myproxyPing	pass
Appliance	
<ul style="list-style-type: none">• rocks-131.sdsc.edu_mdsInfo• rocks-131.sdsc.edu_wsgram• rocks-160.sdsc.edu_mdsInfo• rocks-160.sdsc.edu_wsgram	
ucsd-portal	
rocks-131.sdsc.edu_mdsInfo	pass
rocks-131.sdsc.edu_wsgram	pass
rocks-160.sdsc.edu_mdsInfo	pass
rocks-160.sdsc.edu_wsgram	pass

Discussion: possibilities for expanding monitoring

- What other services should be monitored?
- Expand to more campuses?


Inca Information

- Announcements:
inca-users@sdsc.edu
 - Email:
inca@sdsc.edu
 - Website:
<http://inca.sdsc.edu>
- Supported by:



Related Grid monitoring tools

BIG BROTHER™



Nagios®

**Service Availability
Monitoring**

Inca's primary objective: user-level Grid monitoring